

# Neural Network Analysis

In this project, was built an initial neural network (shallow) used to predict pass through medicines. Was provided a relation of 2 quarter (7 months) of stock with ARIMA prediction & a quarter prediction (2019/03).

In [1]:

```
%matplotlib inline
%config InlineBackend.figure_format = 'retina'

import copy
import progressbar
import numpy as np
import pandas as pd
from datetime import datetime
import matplotlib.pyplot as plt
from matplotlib.pyplot import *
```

In [2]:

```
# Production only
# from google.colab import drive
# drive.mount('/content/drive')
```

## Load and prepare the data

A critical step in working with neural networks is preparing the data correctly. Variables on different scales make it difficult for the network to efficiently learn the correct weights. Below, we've written the code to load and prepare the data. You'll learn more about this soon!

In [3]:

```
arimaset_legend = '(ARIMA)QUETIAPINA'
microset_legend = '(MICROSET)QUETIAPINA'

# Production Path
# arimaset_path = '/content/drive/My Drive/DAF-USP/Preditivo/Colab Notebooks/data/arimaset-quetiapina-alagoas.csv'
# microset_path = '/content/drive/My Drive/DAF-USP/Preditivo/Colab Notebooks/data/microset-quetiapina-alagoas.csv'

arimaset_path = 'datasets/arimaset-quetiapina-alagoas.csv'
microset_path = 'datasets/microset-quetiapina-alagoas.csv'

arimaset_df = pd.read_csv(arimaset_path)
microset_df = pd.read_csv(microset_path)

arimaset_df['DT'] = pd.to_datetime(arimaset_df['DT'], format='%Y%m', errors='ignore')
microset_df['DT'] = pd.to_datetime(microset_df['DT'], format='%Y%m', errors='ignore')
```

In [4]:

```
arimaset_df.head()
```

Out[4]:

	DT	MES	QTD
0	2013-01-01	1	63932
1	2013-01-01	1	116474
2	2013-01-01	1	71167
3	2013-01-01	1	1306
4	2013-01-01	1	62842

In [5]:

```
microset_df.head()
```

Out[5]:

	DT	MES	QTD
0	2011-12-01	12	390
1	2012-01-01	1	1500
2	2012-02-01	2	1875
3	2012-03-01	3	2785
4	2012-04-01	4	3525

In [6]:

```
print('Ajusted QTD (DAF has 3% of Data)')
microset_df.QTD = microset_df.QTD * 33.33
microset_df.head()
```

Ajusted QTD (DAF has 3% of Data)

Out[6]:

	DT	MES	QTD
0	2011-12-01	12	12998.70
1	2012-01-01	1	49995.00
2	2012-02-01	2	62493.75
3	2012-03-01	3	92824.05
4	2012-04-01	4	117488.25

## Checking out the data


This dataset has the number of medicines acquired by Health Ministry for each month of each day from January 1 2017 to November 31 2018. The number of medicines is split between casual and registered, summed up in the `count1` column. You can see the first few rows of the data above.

In [7]:

```
fig, ax = subplots()
arimaset_df[:].plot(x='DT', y=['QTD'], ax=ax)
ax.legend([arimaset_legend])
```

Out[7]:

&lt;matplotlib.legend.Legend at 0x10fd76290&gt;



## DBSet with Outlier


In [8]:

```
fig, ax = subplots()
# microset_df = microset_df[microset_df.DT < datetime(2017,9,1)]
microset_df[:].plot(x='DT', y=['QTD'], ax=ax)

ax.legend([microset_legend])
```

Out[8]:

&lt;matplotlib.legend.Legend at 0x1144aa890&gt;



## Outlier Clear


In [9]:

```
fig, ax = subplots()
microset_df = microset_df[microset_df.DT < datetime(2017,9,1)]
microset_df[:].plot(x='DT', y=['QTD'], ax=ax)

ax.legend([microset_legend])
```

Out[9]:

<matplotlib.legend.Legend at 0x10fd67e10>



## Dummy variables

Here we have some categorical variables like **predicted** and **delivered**. To include these variables is necessary to make it binary dummy variables. This is simple to do with Pandas thanks to `get_dummies()`.

In [10]:

```
# Convert dtmonth to index month ignoring it year
# medicines['dtmonth'] = medicines['dtmonth'].apply(lambda x: int(x.split('-')[1]))

dummy_fields = ['DT']
for each in dummy_fields:
    dummies = pd.get_dummies(microset_df[each], prefix=each, drop_first=False)
    microset = pd.concat([microset_df, dummies], axis=1)

fields_to_drop = ['DT']
data = microset.drop(fields_to_drop, axis=1)
data.head()
```

Out[10]:

			DT_2011-12-01 00:00:00	DT_2012-01-01 00:00:00	DT_2012-02-01 00:00:00	DT_2012-03-01 00:00:00	DT_2012-04-01 00:00:00	DT_2012-05-01 00:00:00	DT_2012-06-01 00:00:00
MES	QTD								
0	12	12998.70	1	0	0	0	0	0	0
1	1	49995.00	0	1	0	0	0	0	0
2	2	62493.75	0	0	1	0	0	0	0
3	3	92824.05	0	0	0	1	0	0	0
4	4	117488.25	0	0	0	0	1	0	0

5 rows x 71 columns

In [11]:

```

dummy_fields = ['DT']
for each in dummy_fields:
    dummies = pd.get_dummies(arimaset_df[each], prefix=each, drop_first=False)
    arimaset = pd.concat([arimaset_df, dummies], axis=1)

fields_to_drop = ['DT']
arima_data = arimaset.drop(fields_to_drop, axis=1)
arima_data.head()

```

Out[11]:

	MES	QTD	DT_2013-01-01 00:00:00	DT_2013-02-01 00:00:00	DT_2013-03-01 00:00:00	DT_2013-04-01 00:00:00	DT_2013-05-01 00:00:00	DT_2013-06-01 00:00:00	DT_2013-07-01 00:00:00	D
0	1	63932	1	0	0	0	0	0	0	
1	1	116474	1	0	0	0	0	0	0	
2	1	71167	1	0	0	0	0	0	0	
3	1	1306	1	0	0	0	0	0	0	
4	1	62842	1	0	0	0	0	0	0	

5 rows × 77 columns

## Scaling target variables

To make training the network easier, we'll standardize each of the continuous variables. That is, we'll shift and scale the variables such that they have zero mean and a standard deviation of 1.

The scaling factors are saved so we can go backwards when we use the network for predictions.

In [12]:

```

quant_features = ['MES', 'QTD']
# Store scalings in a dictionary so we can convert back later
scaled_features = {}
for each in quant_features:
    mean, std = data[each].mean(), data[each].std()
    scaled_features[each] = [mean, std]
    data.loc[:, each] = (data[each] - mean)/std

print('passed')

```

passed

In [13]:

```
quant_features = ['MES', 'QTD']  
# Store scalings in a dictionary so we can convert back later  
scaled_features = {}  
for each in quant_features:  
    mean, std = arima_data[each].mean(), arima_data[each].std()  
    scaled_features[each] = [mean, std]  
    arima_data.loc[:, each] = (arima_data[each] - mean)/std  
  
print('passed')
```

passed

## Splitting the data into training, testing, and validation sets

We'll save the data for the last approximately 2 months to use as a test set after we've trained the network. We'll use this set to make predictions and compare them with the actual number of medicines.

In [14]:

```

# goal_data = daf_df['QTD']
# features_data = daf_df.loc[:, daf_df.columns != 'QTD']

# trainset, testset, validateset = np.split(daf_df.sample(frac=1), [int(.7*len(daf_df)), int(.8*len(daf_df))])
# X_trainset, X_testset, X_validateset = np.split(features_data.sample(frac=1), [int(.7*len(features_data)), int(.8*len(features_data))])
# Y_trainset, Y_testset, Y_validateset = np.split(goal_data.sample(frac=1), [int(.7*len(goal_data)), int(.8*len(goal_data))])

# Save data for approximately the last 2 months
test_data = data[:] # test_data = data[-1*8:]

# Now remove the test data from the data set
data = data[:] #data = data[:-1*8]

# Separate the data into features and targets
target_fields = ['QTD']
features, targets = data.drop(target_fields, axis=1), data[target_fields] # features(inputs) = data except targets
test_features, test_targets = test_data.drop(target_fields, axis=1), test_data[target_fields]
# arima_targets = arimaset_df[(arimaset_df.DT > datetime(2017,1,1)) & (arimaset_df.DT < datetime(2017,9,1))][target_fields]
# arima_targets = arima_targets[-1*8:]
arima_targets = arima_data[-1*8:]

test_data.head()

```

Out[14]:

	MES	QTD	DT_2011-12-01 00:00:00	DT_2012-01-01 00:00:00	DT_2012-02-01 00:00:00	DT_2012-03-01 00:00:00	DT_2012-04-01 00:00:00	DT_2012-05-01 00:00:00	DT_2012-06-01 00:00:00
0	1.629993	-2.608054	1	0	0	0	0	0	0
1	-1.542224	-2.273090	0	1	0	0	0	0	0
2	-1.253841	-2.159926	0	0	1	0	0	0	0
3	-0.965457	-1.885316	0	0	0	1	0	0	0
4	-0.677074	-1.662006	0	0	0	0	1	0	0

5 rows × 71 columns



In [15]:

```
arima_data.head()
```

Out[15]:

	MES	QTD	DT_2013-01-01 00:00:00	DT_2013-02-01 00:00:00	DT_2013-03-01 00:00:00	DT_2013-04-01 00:00:00	DT_2013-05-01 00:00:00	DT_2013-06-01 00:00:00	DT_2013-07-01 00:00:00
0	-1.547998	1.146847	1	0	0	0	0	0	0
1	-1.547998	2.605421	1	0	0	0	0	0	0
2	-1.547998	1.347691	1	0	0	0	0	0	0
3	-1.547998	-0.591661	1	0	0	0	0	0	0
4	-1.547998	1.116588	1	0	0	0	0	0	0

5 rows × 77 columns

We'll split the data into two sets, one for training and one for validating as the network is being trained. Since this is time series data, we'll train on historical data, then try to predict on future data (the validation set).

In [16]:

```
# Hold out the last 8 months or so of the remaining data as a validation set
train_features, train_targets = features[:-2*8], targets[:-2*8]
val_features, val_targets = features[-2*8:], targets[-2*8:]
```

## Time to build the network

Below you'll build your network. We've built out the structure and the backwards pass. **You'll implement the forward pass** through the network. **You'll also set** the hyperparameters: **the learning rate**, **the number of hidden units**, and **the number of training passes**.

The network has two layers: a **hidden** layer and an **output** layer. The **hidden** layer will use the **sigmoid function for activations**. The **output** layer has only one node and is **used for the regression**, the **output** of the node **is the same as the input** of the node. That is, **the activation function is  $f(x) = x$** . A function that takes the input signal and generates an output signal, but takes into account the threshold, is called an activation function. We work through each layer of our network calculating the outputs for each neuron. All of the **outputs from one layer become inputs to the neurons on the next layer**. This process is **called forward propagation**.

We **use the weights to propagate signals forward** from the **input to the output** layers in a neural network. We use the **weights to also propagate error backwards** from the **output back** into the network to update our weights. This is called **backpropagation**.

**Hint:** You'll need the derivative of the output activation function ( $f(x) = x$ ) for the backpropagation implementation. If you aren't familiar with calculus, this function is equivalent to the equation  $y = x$ . What is the slope of that equation? That is the derivative of  $f(x)$ .

Below, you have these tasks:

1. Implement the sigmoid function to use as the activation function. Set `self.activation_function` in `__init__` to your sigmoid function.
2. Implement the forward pass in the `train` method.
3. Implement the backpropagation algorithm in the `train` method, including calculating the output error.
4. Implement the forward pass in the `run` method.

In [17]:

```

class NeuralNetwork(object):
    def __init__(self, input_nodes, hidden_nodes, output_nodes, learning_rate):
        # Set number of nodes in input, hidden and output layers.
        self.input_nodes = input_nodes
        self.hidden_nodes = hidden_nodes
        self.output_nodes = output_nodes

        # Initialize weights
        self.weights_input_to_hidden = np.random.normal(0.0, self.hidden_nodes *
* -0.5,
                                                    (self.hidden_nodes, self
.input_nodes))

        self.weights_hidden_to_output = np.random.normal(0.0, self.output_nodes
** -0.5,
                                                    (self.output_nodes, sel
f.hidden_nodes))
        self.lr = learning_rate

        ##### TODO: Set self.activation_function to your implemented sigmoid func
tion #####
        #
        # Note: in Python, you can define a function with a lambda expression,
        # as shown below.
        self.activation_function = lambda x: 1 / (1 + np.exp(-x))
        self.activation_function_prime = lambda x: x*1

        ### If the lambda code above is not something you're familiar with,
        # You can uncomment out the following three lines and put your
        # implementation there instead.
        #
        # def sigmoid(x):
        #     return 0 # Replace 0 with your sigmoid calculation here
        # self.activation_function = sigmoid

    def train(self, inputs_list, targets_list):
        # Convert inputs list to 2d array
        inputs = np.array(inputs_list, ndmin=2).T
        targets = np.array(targets_list, ndmin=2).T

        ##### Implement the forward pass here #####
        ### Forward pass ###
        # TODO: Hidden layer - Replace these values with your calculations.
        # weights is (n x 56), (a set of weights for each hidden node)
        # hidden inputs linear combination & generating activated output
        hidden_inputs = np.dot(self.weights_input_to_hidden, inputs) # signals
into hidden layer
        hidden_outputs = self.activation_function(hidden_inputs) # signals from
hidden layer

        # TODO: Output layer - Replace these values with your calculations.
        # values on output layer (end nodes), same steps before (linear comb+act
ivation) using out_refs
        final_inputs = np.dot(self.weights_hidden_to_output, hidden_outputs) #
signals into final output layer
        final_outputs = self.activation_function_prime(final_inputs) # the deriv
ate activation function here is f(x)=x, so it is the same

        ##### Implement the backward pass here #####

```

```

#### Backward pass ####

# TODO: Output error - Replace this value with your calculations.
output_errors = targets_list - final_outputs # (y - y^)Output layer error
# error_term: y-y^(chapel)*f'(h), f'(h)=1, so error = error_term

# TODO: Backpropagated error - Replace these values with your calculations.
hidden_errors = output_errors * self.weights_hidden_to_output # errors propagated to the hidden layer
hidden_grad = hidden_outputs * (1 - hidden_outputs) # hidden layer gradients

# TODO: Update the weights - Replace these values with your calculations.
# wi = wi + n(learning-rate)*error-term*xi(inputs) --> error-term
error_term = (hidden_errors.T * hidden_grad)
self.weights_input_to_hidden += self.lr * np.dot(error_term, inputs.T)
# wi = wi + n(learning-rate)*output_errors*outputs_previous_layer(hidden) --> error-term
self.weights_hidden_to_output += self.lr * output_errors * hidden_output

def run(self, inputs_list):
    # Run a forward pass through the network
    inputs = np.array(inputs_list, ndmin=2).T

    #### Implement the forward pass here ####
    # TODO: Hidden layer - replace these values with the appropriate calculations.
    # same from training function, but now using trained weights
    hidden_inputs = np.dot(self.weights_input_to_hidden, inputs)
    hidden_outputs = self.activation_function(hidden_inputs)

    # TODO: Output layer - Replace these values with the appropriate calculations.
    final_inputs = np.dot(self.weights_hidden_to_output, hidden_outputs) # linear combination of weights & hidden output
    final_outputs = self.activation_function_prime(final_inputs) # it will return final_inputs because f'(x) = x

    return final_outputs

```

In [18]:

```

def MSE(y, Y):
    return np.mean((y-Y)**2)

```

## Training the network

Here you'll set the hyperparameters for the network. The strategy here is to find hyperparameters such that the error on the training set is low, but you're not overfitting to the data. If you train the network too long or have too many hidden nodes, it can become overly specific to the training set and will fail to generalize to the validation set. That is, the loss on the validation set will start increasing as the training set loss drops.

You'll also be using a method known as **Stochastic Gradient Descent (SGD) to train the network**. The idea is that for **each training pass**, you **grab a random sample of the data instead** of using the **whole data set**. You use **many more training passes than** with **normal gradient descent**, but each pass is **much faster**. This ends up **training the network more efficiently**. You'll learn more about SGD later.

### Choose the number of epochs

This is the number of times the dataset will pass through the network, each time updating the weights. As the number of epochs increases, the network becomes better and better at predicting the targets in the training set. You'll need to choose enough epochs to train the network well but not too many or you'll be overfitting.

### Choose the learning rate

This scales the size of weight updates. If this is too big, the weights tend to explode and the network fails to fit the data. A good choice to start at is 0.1. If the network has problems fitting the data, try reducing the learning rate. Note that the lower the learning rate, the smaller the steps are in the weight updates and the longer it takes for the neural network to converge.

### Choose the number of hidden nodes

The more hidden nodes you have, the more accurate predictions the model will make. Try a few different numbers and see how it affects the performance. You can look at the losses dictionary for a metric of the network performance. If the number of hidden units is too low, then the model won't have enough space to learn and if it is too high there are too many options for the direction that the learning can take. The trick here is to find the right balance in number of hidden units you choose.

In [19]:

```

import sys

### Set the hyperparameters here ###
epochs = 1600 #16000
learning_rate = 0.000001 #0.000001
hidden_nodes = 70 # Neurons (210=overfit, 70=ok)
output_nodes = 1

N_i = train_features.shape[1]
network = NeuralNetwork(N_i, hidden_nodes, output_nodes, learning_rate)

bar = progressbar.ProgressBar(maxval=20, widgets=[progressbar.Bar('=', '[', ']'),
    ), ' ', progressbar.Percentage()])
# bar.start()
start_time = datetime.now()
losses = {'train':[], 'validation':[]}
for e in range(epochs):
    # Go through a random batch of 128 records from the training data set
    batch = np.random.choice(train_features.index, size=5000)
    for record, target in zip(train_features.loc[batch].values,
        train_targets.loc[batch]['QTD']):
        network.train(record, target)

    # Printing out the training progress
    train_loss = MSE(network.run(train_features), train_targets['QTD'].values)
    val_loss = MSE(network.run(val_features), val_targets['QTD'].values)
    progress = 100 * e/float(epochs)
    #bar.update(progress)
    sys.stdout.write("\rProgress: " + str(progress)[:4] \
        + "% ... Training loss: " + str(train_loss)[:5] \
        + " ... Validation loss: " + str(val_loss)[:5])

    losses['train'].append(train_loss)
    losses['validation'].append(val_loss)
end_time = datetime.now()
# bar.finish()
print('\nSpent Time: ', end_time - start_time)

```


Progress: 99.9% ... Training loss: 0.244 ... Validation loss: 1.977  
 Spent Time: 0:06:07.271490

In [20]:

```
plt.plot(losses['train'], label='Training loss')
plt.plot(losses['validation'], label='Validation loss')
plt.legend()
plt.ylim(top=5)
# plt.ylim(top=20000, bottom=-100000)
```

Out[20]:

(-0.8521190607491803, 5)



## Check out your predictions

Here, use the test data to view how well your network is modeling the data. If something is completely wrong here, make sure each step in your network is implemented correctly.

In [21]:

```
arima_targets.size
```

Out[21]:

616

In [22]:

```
test_targets.size
```

Out[22]:

69

In [23]:

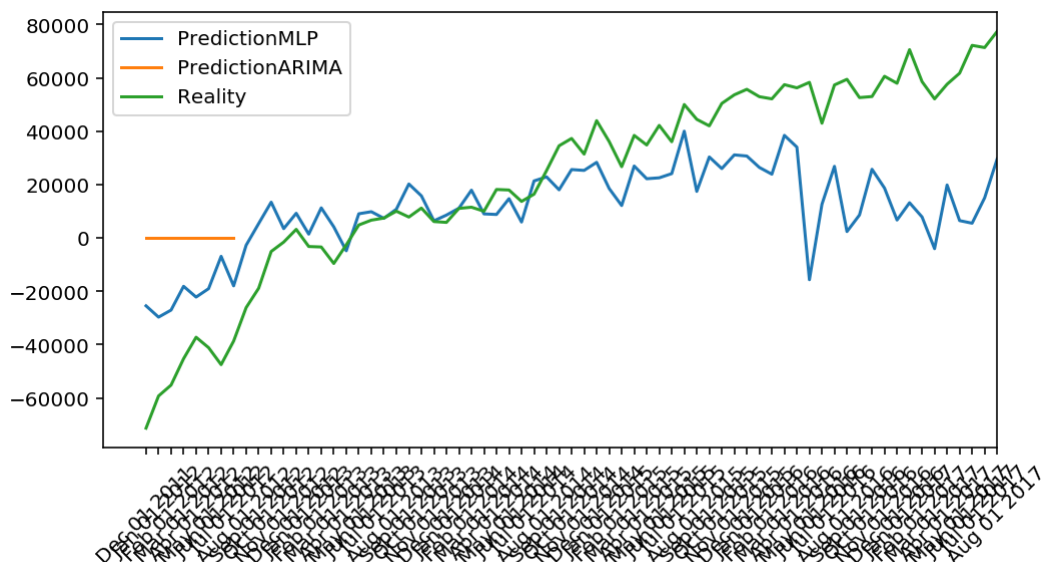
```

fig, ax = plt.subplots(figsize=(8,4))
# plt.yscale('symlog', lenthreshy=1000)

mean, std = scaled_features['QTD']
predictions = network.run(test_features)*std + mean
ax.plot(predictions[0], label='PredictionMLP')
ax.plot((arima_targets['QTD']*std + mean).values, label='PredictionARIMA')
ax.plot((test_targets['QTD']*std + mean).values, label='Reality')
ax.set_xlim(right=len(predictions))
ax.legend()

dates = pd.to_datetime(microset.loc[test_data.index]['DT'])
dates = dates.apply(lambda d: d.strftime('%b %d %Y'))
ax.set_xticks(np.arange(len(dates))[::])
_ = ax.set_xticklabels(dates[:,], rotation=45)

```





In [24]:

```
predictions
```

Out[24]:

```
array([[ -25480.1454082 , -29715.81362803, -27069.76509658,
        -18150.39544995, -22192.53712932, -19032.29778647,
        -6912.76415508, -17985.83485097, -2738.2332511 ,
         5375.57821441,  13394.29766836,   3466.97964678,
         9229.2003198 ,   1398.81141868,  11257.33351463,
         4146.89017245, -4850.61185657,   9035.91023782,
         9816.18480796,   7338.67264741,  10751.57060621,
        20200.48636577,  15762.82567959,   6398.62992495,
         8546.30687993,  11134.24857349,  17892.41286704,
         8994.82698808,   8779.00674425,  14700.84703739,
         5926.8987862 ,  21351.266026 ,  22871.31972267,
        18019.35767269,  25619.60345259,  25295.2202953 ,
        28319.1149652 ,  18525.09381373,  12119.18115262,
        26947.72168778,  22176.91048834,  22516.66625684,
        24091.98820649,  40002.49248333,  17431.14446193,
        30335.33982021,  25967.32624782,  31117.20201218,
        30683.96151426,  26412.29372511,  23877.24829512,
        38438.15463492,  34029.19996723, -15700.43459979,
        12576.88360093,  26824.93587157,   2366.86530026,
         8578.18040503,  25727.42668172,  18676.82759245,
         6688.00171706,  13181.07610758,   7842.6164634 ,
        -4051.04245336,  19829.82060053,   6477.78875475,
         5500.88593265,  15004.60080876,  29576.03999808]])
```

In [25]:

```
test_features
```

Out[25]:

	MES	DT_2011- 12-01 00:00:00	DT_2012- 01-01 00:00:00	DT_2012- 02-01 00:00:00	DT_2012- 03-01 00:00:00	DT_2012- 04-01 00:00:00	DT_2012- 05-01 00:00:00	DT_2012- 06-01 00:00:00	DT_2012- 07-01 00:00:00
0	1.629993	1	0	0	0	0	0	0	0
1	-1.542224	0	1	0	0	0	0	0	0
2	-1.253841	0	0	1	0	0	0	0	0
3	-0.965457	0	0	0	1	0	0	0	0
4	-0.677074	0	0	0	0	1	0	0	0
...	...	...	...	...	...	...	...	...	...
64	-0.677074	0	0	0	0	0	0	0	0
65	-0.388691	0	0	0	0	0	0	0	0
66	-0.100307	0	0	0	0	0	0	0	0
67	0.188076	0	0	0	0	0	0	0	0
68	0.476459	0	0	0	0	0	0	0	0

69 rows × 10 columns

In [26]:

```
scaled_features
```

Out[26]:

```
{'MES': [6.340425531914893, 3.4498918614978353],  
  'QTD': [22619.313829787236, 36022.85356023201]}
```

In [27]:

```
arima_targets['QTD']
```

Out[27]:

```
1120    -0.627916  
1121    -0.627916  
1122    -0.627916  
1123    -0.627916  
1124    -0.627916  
1125    -0.627916  
1126    -0.627916  
1127    -0.627916  
Name: QTD, dtype: float64
```

In [28]:

```
test_targets['QTD']
```

Out[28]:

```
0    -2.608054  
1    -2.273090  
2    -2.159926  
3    -1.885316  
4    -1.662006  
...  
64    0.971234  
65    1.085605  
66    1.375605  
67    1.352671  
68    1.520455  
Name: QTD, Length: 69, dtype: float64
```

In [29]:

```
(arima_targets['QTD']*std + mean).values
```

Out[29]:

```
array([0., 0., 0., 0., 0., 0., 0., 0.])
```

In [30]:

```
(test_targets['QTD']*std + mean).values
```

Out[30]:

```
array([-71330.22309233, -59263.85933226, -55187.38508899, -45295.140
92532,
      -37250.8984186 , -41164.31369214, -47523.61351164, -38674.946
75425,
      -26130.27668296, -18868.71723095, -5095.6695877 , -1573.595
84151,
       3198.59667261, -3236.79733277, -3410.72690048, -9596.097
15227,
      -2606.30264981,  4840.0569679 ,  6644.57623292,  7427.259
28763,
      10014.46160735,  7807.730217 , 11177.61559143,  6122.787
52978,
       5807.5401883 , 11047.16841565, 11503.73353089,  9981.849
81341,
      18145.66889793, 17949.99813425, 13666.98252932, 16384.632
02484,
      25244.16938021, 34527.66005688, 37277.92134634, 31407.798
43603,
      43952.46850732, 36125.63796024, 26679.08831384, 38419.334
13445,
      34821.16620239, 42169.69043826, 36038.67317638, 49996.520
98534,
      44441.64541651, 41995.76087055, 50398.73311067, 53627.300
71134,
      55692.71432793, 52931.58244049, 52116.28759184, 57442.880
60304,
      56247.11482502, 58301.65784362, 42984.98528691, 57323.304
02524,
      59443.07063174, 52594.59390305, 52985.9354304 , 60584.483
41985,
      57964.66930618, 70563.69236738, 58573.42279317, 52105.416
99385,
      57605.93957277, 61725.89620797, 72172.54086872, 71346.375
42208,
      77390.4279001 ])
```

In [31]:

predictions[0]

Out[31]:

```
array([-25480.1454082 , -29715.81362803, -27069.76509658, -18150.395
44995,
      -22192.53712932, -19032.29778647,  -6912.76415508, -17985.834
85097,
      -2738.2332511  ,   5375.57821441,  13394.29766836,   3466.979
64678,
       9229.2003198  ,   1398.81141868,  11257.33351463,   4146.890
17245,
      -4850.61185657,   9035.91023782,   9816.18480796,   7338.672
64741,
      10751.57060621,  20200.48636577,  15762.82567959,   6398.629
92495,
       8546.30687993,  11134.24857349,  17892.41286704,   8994.826
98808,
       8779.00674425,  14700.84703739,   5926.8987862  ,  21351.266
026  ,
      22871.31972267,  18019.35767269,  25619.60345259,  25295.220
2953  ,
      28319.1149652  ,  18525.09381373,  12119.18115262,  26947.721
68778,
      22176.91048834,  22516.66625684,  24091.98820649,  40002.492
48333,
      17431.14446193,  30335.33982021,  25967.32624782,  31117.202
01218,
      30683.96151426,  26412.29372511,  23877.24829512,  38438.154
63492,
      34029.19996723, -15700.43459979,  12576.88360093,  26824.935
87157,
       2366.86530026,   8578.18040503,  25727.42668172,  18676.827
59245,
       6688.00171706,  13181.07610758,   7842.6164634  , -4051.042
45336,
      19829.82060053,   6477.78875475,   5500.88593265,  15004.600
80876,
      29576.03999808])
```

## Result:


Considering 60% of safety range the **BNAFARIUNS** predicted 395,2(247.557771061,6); 377,6(236.284215051,6). Meaning just 7 extra medicines (from 388 (real) to 395 (predicted)) for for 2018-10 and 101 medicines for 2018-11. **Question:** Was it model prediction a good one?

**Note:** The time frame was too short, creating fictional data for more months increased the assertivity.

## Answer

DAF Resp:

## Sample with better data



## Training with better data (loss decay)

